

SHUBHAM RAVINDRA KOSHTI

3159527490 • srkoshti98@gmail.com • <https://www.linkedin.com/in/shubham-koshti> • <https://www.github.com/shubhamkoshti71>

SUMMARY

Data Engineer with 4+ years of experience specializing in unifying fragmented, multi-source data into governed, real-time platforms that power AI-driven products and operational decision-making. Experienced across Snowflake, BigQuery, Databricks, and Redshift, building API-based ingestion pipelines, single-source-of-truth data models, and the real-time applications that run on top of them. Adept at owning data end-to-end in fast-moving, AI-native environments, from third-party system integration and privacy-sensitive data handling to shipping quickly with AI-assisted development.

WORK EXPERIENCE

Data Analytics Engineer | Quantegy Analytics Inc. | Brooklyn, NY **March 2025 – Present**

- Unified six disparate source systems into a three-layer medallion architecture in Snowflake, creating a single governed source of truth (Core Data Model plus 100+ downstream tables) serving Finance, Operations, Compliance, and Scheduling teams.
- Built automated, multi-source ELT pipelines (Redshift to Snowflake via Dagster) with incremental loading and idempotent, integrity-guaranteed runs, transforming raw data into governed models at 50M+ records/day and cutting dashboard load times by 67%.
- Developed a real-time Streamlit platform used by 120+ operational staff across 4 markets — with bidirectional Snowflake writes, daily ingestion, and live history tracking — turning the governed data layer into a live operational control center.
- Leveraged AI-assisted development (Claude Code) to translate business requirements into production data applications in hours instead of days, while maintaining code quality, RBAC, and real-time data architecture across a growing platform.

Analytics Engineer | iConsult Collaborative | Syracuse, NY **August 2023 – December 2024**

- Built batch ETL pipelines ingesting historical and daily data from multiple sources using Apache Airflow DAGs and Databricks workflows, reducing manual data collection time by 75%.
- Designed a dimensional data warehouse in Databricks using Kimball modeling, cutting manual report-generation time by 66% and accelerating downstream decision-making.
- Implemented data validation and standardization processes to ensure accuracy and consistency of data feeding downstream reporting and analytics.
- Partnered with cross-functional and executive stakeholders to identify reliable data sources and scale the automated solution to 4 additional regions, saving an estimated 120 person-hours per quarter.

Data Analyst, Marketing Strategy | Infosys Ltd. | Pune, India **June 2020 – April 2022**

- Performed customer segmentation and purchasing-behavior analysis using SQL and Python to surface targeted, revenue-driving insights across the Singapore region.
- Built automated reporting workflows to monitor customer-behavior trends and support data-driven decisions.
- Analyzed multi-stage customer order-journey data from Amazon Redshift to identify delays and redundancies and improve operational efficiency.

PROJECTS

FEMA Disaster Analysis: End-to-End Workflow Automation **September 2024 – December 2024**

- Designed Apache Airflow DAGs to automate recurring ingestion of data from external third-party APIs into cloud storage, eliminating manual intervention.
- Built a scalable Airflow + Databricks pipeline processing ~1M+ records across multiple datasets, using Unity Catalog, PySpark, and SparkSQL for consistent, high-quality data.
- Established a real-time connection between Tableau and Databricks to deliver interactive dashboards on the processed data.

Predictive Employee Attrition Analysis & Modeling **April 2024 – June 2024**

- Built predictive models (Logistic Regression, Random Forest) achieving 96% accuracy to identify key behavioral drivers from historical data.
- Translated model output into actionable recommendations projecting a 15% reduction in attrition.

EDUCATION

Syracuse University | Master of Science, Information Systems (Data Science) **May 2024**

Pune University (I2IT) | Bachelor of Engineering, Electronics & Telecommunications **May 2020**

TECHNICAL SKILLS

Languages: Python, SQL, R,, C++

Data Warehouses & Lakes: Snowflake, BigQuery, Databricks, Redshift

Pipelines & Orchestration: Apache Airflow, Dagster, ELT / incremental loading, dbt-style modeling, PySpark, SparkSQL

Cloud & Infrastructure: AWS, Azure, GCP, Kubernetes, Azure Blob Storage, Unity Catalog

Data Apps & BI: Streamlit, Tableau, Dash / Plotly, Power BI

Modeling & Practices: Medallion & Kimball data modeling, API / third-party integration, real-time & batch processing, data quality & validation, RBAC & governance, AI-assisted development (Claude Code)